Final sweep, host kernel: bandwidth relative to the defaults, strided_transfer (4-byte rows, M = 2 and 4)

Each cell: top row AdaptiveCpp, bottom row DPC++; left to right 64Ki, 1Mi, 4Mi, 16Mi rows per copy. Median of 3 runs, geometric mean over M.
Reference: the anchor of the same job (same node, minutes apart) = host-kernel defaults: 2 queues x 4 MiB, MAX_ELEM_SIZE 16, 1 elem/item, WG 32, SDMA off.
Blue faster, red slower, 0.5x to 2x. Radeon 8060S AdaptiveCpp H>D drifts with time: read with care.

